Custom View Settings

Question #11 *Topic 1*

Which of the following are characteristics of Amazon VPC subnets? (Choose two.)

A. Each subnet maps to a single Availability Zone

B. A CIDR block mask of /25 is the smallest range supported

C. Instances in a private subnet can communicate with the internet only if they have an Elastic IP.

D. By default, all subnets can route between each other, whether they are private or public

E. V Each subnet spans at least 2 Availability zones to provide a high-availability environment

Correct Answer: *AD*
ג€Each subnet must reside entirely within one Availability Zone and cannot span zones.ג€
ג€Every subnet that you create is automatically associated with the main route table for the VPC.ג€
Reference:
http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Subnets.html

Question #12 *Topic 1*

You are creating an Auto Scaling group whose Instances need to insert a custom metric into CloudWatch.
Which method would be the best way to authenticate your CloudWatch PUT request?

A. Create an IAM role with the Put MetricData permission and modify the Auto Scaling launch configuration to launch instances in that role

B. Create an IAM user with the PutMetricData permission and modify the Auto Scaling launch configuration to inject the userscredentials into the instance User Data

C. Modify the appropriate Cloud Watch metric policies to allow the Put MetricData permission to instances from the Auto Scaling group

D. Create an IAM user with the PutMetricData permission and put the credentials in a private repository and have applications on the server pull the credentials as needed

Correct Answer: *A*
Creates an IAM role is always the best practice to give permissions to EC2 instances in order to interact with other AWS services

Question #13 *Topic 1*

When an EC2 instance that is backed by an S3-based AMI Is terminated, what happens to the data on me root volume?

A. Data is automatically saved as an E8S volume.

B. Data is automatically saved as an ESS snapshot.

C. Data is automatically deleted.

D. Data is unavailable until the instance is restarted.

Correct Answer: *C*

We recommend that you use AMIs backed by Amazon EBS, because they launch faster and use persistent storage.

Reference:

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/RootDeviceStorage.html#choose-an-ami-by-root-device

Community vote distribution

C (100%)

Question #14 *Topic 1*

You have a web application leveraging an Elastic Load Balancer (ELB) In front of the web servers deployed using an Auto Scaling Group Your database is running on Relational Database Service (RDS) The application serves out technical articles and responses to them in general there are more views of an article than there are responses to the article. On occasion, an article on the site becomes extremely popular resulting in significant traffic Increases that causes the site to go down.

What could you do to help alleviate the pressure on the infrastructure while maintaining availability during these events? (Choose three.)

A. Leverage CloudFront for the delivery of the articles.

B. Add RDS read-replicas for the read traffic going to your relational database

C. Leverage ElastiCache for caching the most frequently used data.

D. Use SOS to queue up the requests for the technical posts and deliver them out of the queue.

E. Use Route53 health checks to fail over to an S3 bucket for an error page.

Correct Answer: *ABC*

Question #15 *Topic 1*

The majority of your Infrastructure is on premises and you have a small footprint on AWS Your company has decided to roll out a new application that is heavily dependent on low latency connectivity to LOAP for authentication Your security policy requires minimal changes to the company's existing application user management processes.
What option would you implement to successfully launch this application1?

A. Create a second, independent LOAP server in AWS for your application to use for authentication

B. Establish a VPN connection so your applications can authenticate against your existing on-premises LDAP servers

C. Establish a VPN connection between your data center and AWS create a LDAP replica on AWS and configure your application to use the LDAP replica for authentication

D. Create a second LDAP domain on AWS establish a VPN connection to establish a trust relationship between your new and existing domains and use the new domain for authentication

Correct Answer: *C*
Create read replica(RODC) of main LDAP server so that LDAP read replica or RODC can authenticate with application locally.
Creating new domain and trust relationship would require lot of work and changes in exiting ldap configuration so D cannot be answer here.

Community vote distribution
B (100%)

Question #16 *Topic 1*

You need to design a VPC for a web-application consisting of an Elastic Load Balancer (ELB). a fleet of web/application servers, and an RDS database. The entire Infrastructure must be distributed over 2 availability zones.
Which VPC configuration works while assuring the database is not available from the Internet?

A. One public subnet for ELB one public subnet for the web-servers, and one private subnet for the database

B. One public subnet for ELB two private subnets for the web-servers, two private subnets for RDS

C. Two public subnets for ELB two private subnets for the web-servers and two private subnets for RDS

D. Two public subnets for ELB two public subnets for the web-servers, and two public subnets for RDS

Correct Answer: *C*
While using ELB for web applications, ensure that you place all other EC2 instances in private subnets wherever possible. Except where there is an explicit requirement for instances requiring outside world access and Elastic IP attached, place all the instances in private subnets only. In the Amazon VPC environment, only ELBs must be in the public subnet as secure practice.
You will need to select a Subnet for each Availability Zone where you wish traffic to be routed by your load balancer. If you have instances in only one Availability
Zone, please select at least two Subnets in different Availability Zones to provide higher availability for your load balance

Question #17 *Topic 1*

An application that you are managing has EC2 instances & Dynamo OB tables deployed to several AWS Regions in order to monitor the performance of the application globally, you would like to see two graphs:
1) Avg CPU Utilization across all EC2 instances
2) Number of Throttled Requests for all DynamoDB tables.
How can you accomplish this?

A. Tag your resources with the application name, and select the tag name as the dimension in the Cloudwatch Management console to view the respective graphs

B. Use the Cloud Watch CLI tools to pull the respective metrics from each regional endpoint Aggregate the data offline & store it for graphing in CloudWatch.

C. Add SNMP traps to each instance and DynamoDB table Leverage a central monitoring server to capture data from each instance and table Put the aggregate data into Cloud Watch for graphing.

D. Add a CloudWatch agent to each instance and attach one to each DynamoDB table. When configuring the agent set the appropriate application name & view the graphs in CloudWatch.

Correct Answer: *B*
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Tools.CLI.html

Question #18 *Topic 1*

When assessing an organization s use of AWS API access credentials which of the following three credentials should be evaluated? (Choose three.)

A. Key pairs

B. Console passwords

C. Access keys

D. Signing certificates

E. Security Group memberships

Correct Answer: *BCD*
AWS provides a number of authentication mechanisms including a console, account IDs and secret keys, X.509 certificates, and MFA devices to control access to
AWS APIs. Console authentication is the most appropriate for administrative or manual activities, account IDs and secret keys for accessing REST-based interfaces or tools, and X.509 certificates for SOAP-based interfaces and tools.
Your organization should consider the circumstances under which it will leverage access keys, x.509certificates, console passwords, or MFA devices

Question #19 *Topic 1*

You have a Linux EC2 web server instance running inside a VPC The instance is In a public subnet and has an EIP associated with it so you can connect to It over the Internet via HTTP or SSH The instance was also fully accessible when you last logged in via SSH. and was also serving web requests on port 80.
Now you are not able to SSH into the host nor does it respond to web requests on port 80 that were working fine last time you checked You have double-checked that all networking configuration parameters (security groups route tables. IGW'EIP. NACLs etc) are properly configured {and you haven't made any changes to those anyway since you were last able to reach the Instance). You look at the EC2 console and notice that system status check shows "impaired."
Which should be your next step in troubleshooting and attempting to get the instance back to a healthy state so that you can log in again?

A. Stop and start the instance so that it will be able to be redeployed on a healthy host system that most likely will fix the "impaired" system status

B. Reboot your instance so that the operating system will have a chance to boot in a clean healthy state that most likely will fix the 'impaired" system status

C. Add another dynamic private IP address to me instance and try to connect via mat new path, since the networking stack of the OS may be locked up causing the ג€impairedג€ system status.

D. Add another Elastic Network Interface to the instance and try to connect via that new path since the networking stack of the OS may be locked up causing the "impaired" system status

E. un-map and then re-map the EIP to the instance, since the IGWVNAT gateway may not be working properly, causing the "impaired" system status

Correct Answer: *A*

Question #20 *Topic 1*

What is a placement group?

A. A collection of Auto Scaling groups in the same Region

B. Feature that enables EC2 instances to interact with each other via nigh bandwidth, low latency connections

C. A collection of Elastic Load Balancers in the same Region or Availability Zone

D. A collection of authorized Cloud Front edge locations for a distribution

Correct Answer: *B*

Previous Questions Next Questions